

Error handling in modern C++



Mikel Irazabal

12/11/2019

Content

1 Introduction

2 Contracts

3 Error Codes

4 Exceptions

5 Expected

6 Throwing by value

7 Conclusion

8 Questions

Content

1 Introduction

2 Contracts

3 Error Codes

4 Exceptions

5 Expected

6 Throwing by value

7 Conclusion

8 Questions

Introduction

- error: lat. *error*, -ōris.
- m. Acción desacertada o equivocada. - [RAE online]
- error: “an act that ... fails to achieve what should be done.” — [Merriam-Webster]
- Error: a function that couldn’t do what it advertised (i.e., its preconditions were met, but it could not achieve its successful-return postconditions, and the calling code can recover).
- Bug: state of the program not foreseen by the designer.
- A bug or abstract machine corruption is not an error (both are not programmatically recoverable)
- This includes out-of-bound access and null dereference.
- A programmer is needed to recover from such a failure.

Content

1 Introduction

2 Contracts

3 Error Codes

4 Exceptions

5 Expected

6 Throwing by value

7 Conclusion

8 Questions

Contracts

- Derive from Hoare logic.
- Term coined by Bertrand Meyer and made it available in Eiffel.
- Idea of having assertions on steriods.
- Define Pre-condition (expect), assert and post-condition (ensure) (e.g., invariants were left out)
- Define *audit* (expensive operation) and *axiom* (mostly for static analyzers)



Figure: Sign it!

Contracts

- “With contracts, 90-something% of the typical uses of exceptions in .NET and Java became preconditions. All of the ArgumentNullException, ArgumentOutOfRangeException, and related types and, more importantly, the manual checks and throws were gone.” — [Duffy 2016]



Figure: Midori OS

Contracts

```
1
2 void f(int x, int y)
3     [[expects: x>0]]
4     [[expects: y!=0]]
5     [[ensures result: result > x+y]];
6
7
8 bool positive(const int * p) [[expects: p!=nullptr]] {
9     return *p > 0;
10 }
11
12 bool g(int * p) [[expects: positive(p)]];
```

Listing 1: Ex. from document P0542R5

Contracts

- Contracts didn't make it for C++20 :-(. Possibly for C++26.
- In any case it has been proof that restricting the sets that are passed to the functions reduces the number of bugs (e.g., prefer passing by const rather than by *, or don't pass int if the number cannot be negative).

Content

1 Introduction

2 Contracts

3 Error Codes

4 Exceptions

5 Expected

6 Throwing by value

7 Conclusion

8 Questions

Error Codes

- Available in C.
- Different approaches can be applied.
 - Global variable (e.g., errno)
 - Return value (e.g., open)
 - Singletons (e.g.,)



Figure: From the ancient times...

Error Codes

- Set errno to 0
- Check if the value changed

```
1
2 int main()
3 {
4     // set errno to 0
5     errno = 0;
6     id_t calling_process = 0;
7     int rc = getpriority(PRIO_PROCESS, calling_process);
8     if(errno != 0)
9         printf("Error getting the priority. Message: %s \n",
10            strerror(errno));
11
12     return 0;
}
```

Listing 2: errno example 1

Error Codes

- Check for return value and then errno

```
1 int main()
2 {
3     const char* file_path= "/tmp/foo.txt";
4     int fd = open(file_path, O_APPEND );
5     if(fd < 0)
6         printf("Error opening the file %s with the message: %s \n",
7                file_path, strerror(errno) );
8     return 0;
9 }
10
```

Listing 3: errno example 1

Error Codes

- Check for return value. The used value to indicate an error, may also be a valid value (e.g., 0 value is used to mark an error).

```
1
2 const char *str1 = "42";
3 const char *str2 = "3.14159";
4 const char *str3 = "BcnCppMeetup";
5 const char *str4 = "0";
6
7 int num1 = std::atoi(str1);
8 int num2 = std::atoi(str2);
9 int num3 = std::atoi(str3); \\ Error. num3 = 0.
10 int num4 = std::atoi(str4); \\ Error or valid input?
11
12 std::cout << "std::atoi(\"" << str1 << "\") is " << num1 << '\n';
13 std::cout << "std::atoi(\"" << str2 << "\") is " << num2 << '\n';
14 std::cout << "std::atoi(\"" << str3 << "\") is " << num3 << '\n';
15 std::cout << "std::atoi(\"" << str4 << "\") is " << num4 << '\n';
```

Listing 4: atoi

- Check if the return value is not equal to zero. errno is not changed on failure!

```
1 const size_t SIMD_BIT_ALIGN = 512;
2 const size_t size = 1000;
3 void* ptr;
4 int rc = posix_memalign( &ptr, SIMD_BIT_ALIGN, size);
5 if(rc != 0){
6     switch(rc): // note that errno doesn't change
7         case EINVAL:{
8             printf("The alignment argument was not a power of two, or
9 was not a multiple of sizeof(void *)\n");
10            break;
11        }
12        case ENOMEM:{
13            printf("There was insufficient memory to fulfill the
14 allocation request\n" );
15            break;
16        }
17        default:
18            assert(0 != 0 && "Flow cannot get here..")
```

Listing 5: posix_memalign example

Error codes

- Possible memory leaks...

```
1 int main()
2 {
3     int* p = malloc(1000*sizeof(int));
4     if(p == NULL){}
5         printf("Error allocating memory");
6         return 0;
7     }
8
9     // very important code
10
11    free(p); // can we guarantee that all the paths pass through
12    the free?
13    return 0;
14 }
```

Listing 6: malloc example

Error codes

- Maybe a valid idiom in C for resource management, use RAII in C++

```
1 int* p = malloc(1000*sizeof(int));
2 if(p == NULL){
3     goto exit;
4 }
5 if(!foo())
6     goto free_func;
7 if(!bar())
8     goto free_func;
9 if(!baz())
10    goto free_func;
11
12 free(p);
13 return 1;
14 free_func:
15     free(p);
16 exit:
17     return 0;
```

Listing 7: malloc example

Error Codes

- And what about singletons for managing errors!!!!
- Composition of errors is not trivial

```
1 int main()
2 {
3     GLenum rc = glGetError();
4     if(rc != GL_NO_ERROR){
5         switch (rc):
6             case GL_INVALID_ENUM: {
7                 printf("Invalid enum passed!");
8                 break;
9             }
10            /*
11            .
12            */
13        }
14        // more important code
15    return 0;
16 }
```

Listing 8: OpenGL example

Error Codes

- Implementation

- `errno` and `perror` are thread safe according to the C11 standard
- `strerror` is not thread safe. Use `strerror_r` as alternative
- Out-parameters enforce a memory layout which is not optimizer friendly, basically the return value that it is stored in a register unusable for normal flow code

- Conceptual

- A lot of noise added
- Separating error path from normal flow is difficult
- Tedious to maintain consistently such an approach
- Postponing error-handling requires the programmer to take into account the error-code through the call-graph.
- Basically if an error is detected, a task cancellation should happen.
- Visible flow

Error Codes

- There was at least one continuous bug in the previous slides, can someone point it out?



Figure: Ouchhhhhhhh!

Error Codes

- Yes, printf also returns a value to inform the caller about the success or failure... Easy to forget checking for all the return values.
- C++17 added an attribute to generate a warning if the caller ignores the return value. `[[nodiscard]]`
- If a function declared `nodiscard` or a function returning an enumeration or class declared `nodiscard` by value is called from a discarded-value expression other than a cast to void, the compiler is encouraged to issue a warning.

Error Codes

- Since C++11 we can integrate the error codes into the language.
- It is composed of an integer and a pointer to the domain where the error happened.

```
1 struct error_code
2 {
3     private:
4         int             _M_value // error number;
5         const error_category* _M_cat // domain where the error was
6         originated;
7 }
```

Listing 9: std::error_code

- `std::error_category` is intended to be used with the typical OOP.

```
1 class error_category
2 {
3     public:
4         constexpr error_category() noexcept = default;
5         virtual ~error_category();
6         error_category(const error_category&) = delete;
7         error_category& operator=(const error_category&) = delete;
8
9         // category name
10        virtual const char* name() const noexcept = 0;
11        // error message
12        virtual string message(int) const = 0;
13        // maps error_code to error_condition
14        virtual error_condition default_error_condition(int __i)
15        const noexcept;
16
17        virtual bool equivalent(int __i, const error_condition& __cond
18        ) const noexcept;
```

Listing 10: `std::error_category`

Error Codes

- Implementation

```
1 // Don't do this!! It exists std::errc !!
2
3 enum class socket_err_c {
4     no_error = 0,      // 0 should not represent an error
5     invalid_addr = 1,
6     unknown_proto = 2,
7     invalid_flags = 3,
8     process_fd_limit = 4,
9     system_fd_limit = 5,
10    insufficient_memory = 6,
11    proto_not_supported = 7,
12 };
13 // Yes, we have to pollute the std namespace
14 namespace std
15 {
16     template <> struct is_error_code_enum<socket_err_c> : true_type
17     { };
18 }
19
```

Listing 11: std::error_category

Error Codes

```
1  class socket_err_c_category : public std::error_category {
2  public:
3      // Return a short descriptive name for the category
4      const char *name() const noexcept final { return "Socket
5      error defined by you!"; }
6      std::string message(int c) const final
7      {
8          switch (static_cast<socket_err_c>(c)) {
9              case socket_err_c::no_error:
10                 return "conversion successful";
11             case socket_err_c::invalid_addr:
12                 return "converting empty string";
13             case socket_err_c::unknown_proto:
14                 return "unknown proto";
15             case socket_err_c::invalid_flags:
16                 return "invalid flags";
17             case socket_err_c:: process_fd_limit:
18                 return "process file descriptors limit reached";
19             default:
20                 return std::to_string(c) + std::string(" code unknown");
21         }
22     };
}
```

Error Codes

```
1 // Declare a singleton
2 socket_err_c_category& get_socket_err_c_category()
3 {
4     static socket_err_c_category ec;
5     return ec;
6 }
7 // Overload the global make_error_code() free function with our
8     custom enum.
9 inline std::error_code make_error_code(socket_err_c ec)
10 {
11     return {static_cast<int>(ec), get_socket_err_c_category()};
12 }
```

Listing 13: std::error_category

Error Codes

```
1 // function call
2 std::error_code open_socket()
3 {
4     // try to open the socket without success...
5     return make_error_code(socket_errc::unknown_proto) ;
6 }
7
8
9 int main()
10 {
11     auto ec = open_socket();
12     assert(ec == socket_errc::no_error);
13     return 0;
14 }
```

Listing 14: std::error_category

Error Codes

- It was adopted from the boost::error_code
- It is supposed to be used at some point by the Networking library
- Nowadays used at the std::filesystem
 - `bool remove(const std::filesystem::path p);`
 - `bool remove(const std::filesystem::path p, std::error_code ec) noexcept;`

Content

1 Introduction

2 Contracts

3 Error Codes

4 Exceptions

5 Expected

6 Throwing by value

7 Conclusion

8 Questions

Exceptions or interrupts

- Exceptions may happen in different domains
 - Hardware exceptions: IEEE 754 floating point exception ex. dividing by zero
 - OS signals. The OS may inform processes via IPC about different behaviours. A handler is registered and the OS can call it. ex. `int$0x80`, SIGKILL.
 - Software exceptions: The term exception is typically used in a specific sense to denote a data structure storing information about an exceptional condition

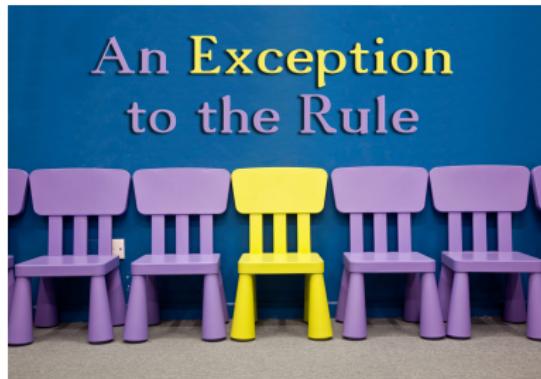


Figure: Exceptions are exceptional!

Exceptions

- In the 80s and with the hope of improving the complexity of canceling errors, a new paradigm arrived in C++ : exceptions
- It has to be noted that C++ can only detect failures through exceptions at implicitly declared member functions (e.g., constructor, destructor...).
- One of the first papers, was written in 89 by Andrew Koenig and Bjarne Stroustrup and they describe how exceptions should be implemented. They discard the resumption ability: *termination is preferred over resumption; this is not a matter of opinion but a matter of years of experience. Resumption is seductive, but not valid.* Bjarne Stroustrup
- A failure can happen in the library but the library doesn't know how to handle it, but the customer of that library does. Ex. `bad_alloc`

Exceptions

```
1 int foo(int size) {
2     void *ptr = malloc(size);
3     if (ptr == 0)
4         throw NoMoreMemory();
5 }
6 int bar(int elem) {
7     Resource res;
8     auto result = foo(2 * (res+elem));
9     /* Do something else */
10 }
11 int main() {
12     try {
13         for (auto i = 0; i < 100; ++i)
14             bar(i);
15     } catch (const NoMoreMemory& excep) {
16         std::cerr << "Out of memory " << '\n';
17     } catch (...) {/* Report other problems. */}
18 }
```

Listing 15: Exceptions usage

Basic guarantee, Strong guarantee and noexcept guarantee

- Good at centralizing error handling mechanism
- They propagate the error up through the stack. Intermediate functions do not have to add code to manage the exception and destructors are called as expected.
- The exception is identified by a specific entity and not by a special return value, and therefore, exceptions liberate the programmer from the hassle of the failure call-graph. Exception code path can be clearly identified.
- There are two basic implementations of exceptions
 - List based (gcc's flag `-RTS=sjlj` in Ada)
 - Table based (gcc's flag `-RTS=zcx` in Ada)

List based exceptions

- Typically, based on setjmp/longjmp functions
- The setjmp function saves the actual context in a jmp_buf structure
- The longjmp is used to perform a non-local goto.
- The try/catch block is replaced with setjmp call and the throw with longjmp.
- A linked-list of jmp_buf structure represents the deepness in the stack.
- Problem: Destruction of the objects. A linked-list with the objects to be destroyed needs to be maintained.

List based exceptions

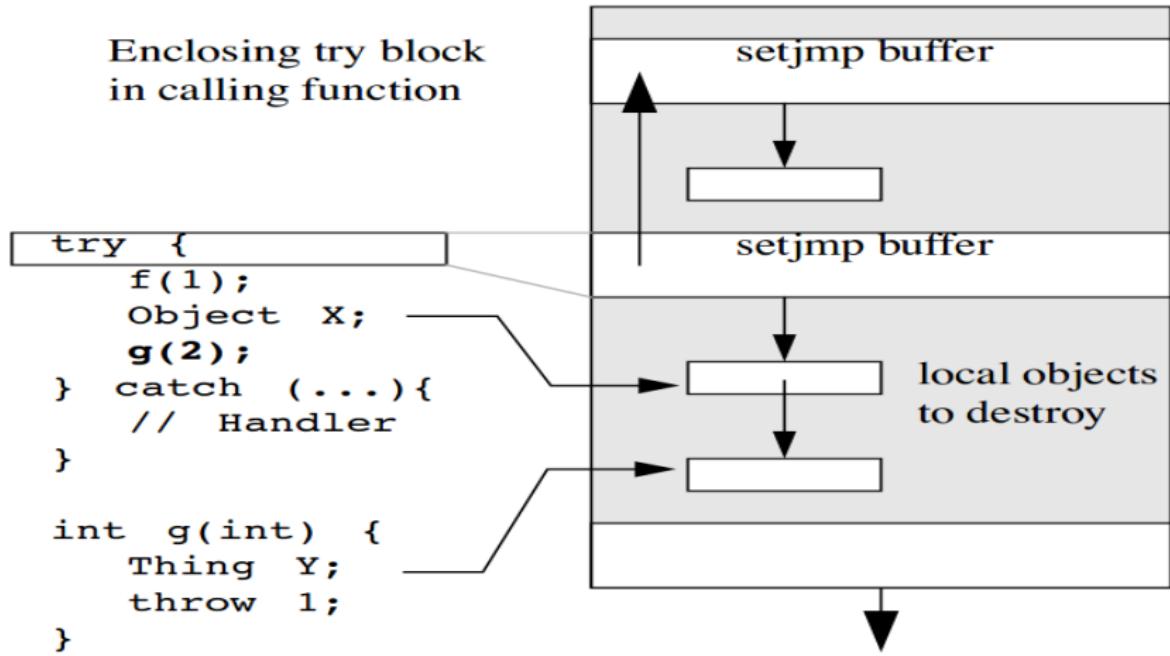


Figure: Setjmp exceptions implementation.

List based exceptions

- Additionally, performance drawbacks are significant
 - The `setjmp` function must be called at every beginning of all the functions, and the `jmp_buf` structure must be maintained.
 - Maintain a list of the object that we may need to destroy
 - Save all the variables values outside the `try` block. This can be a significance overhead, in architectures with large number of registers. Ex. RISC processors.

Table-driven exceptions

- A table is used to map the value of the program counter (PC) at the point where the exception is thrown.
- An action table is used to perform various operations (e.g., invoking the destructors, adjusting the stack or matching the exception thrown with the exception handler, returning the address of the handler).
- E.g., Call the destructor for object X at stack offset Y.
- Control is transferred to the handler once the PC is refreshed.

Table-driven exceptions

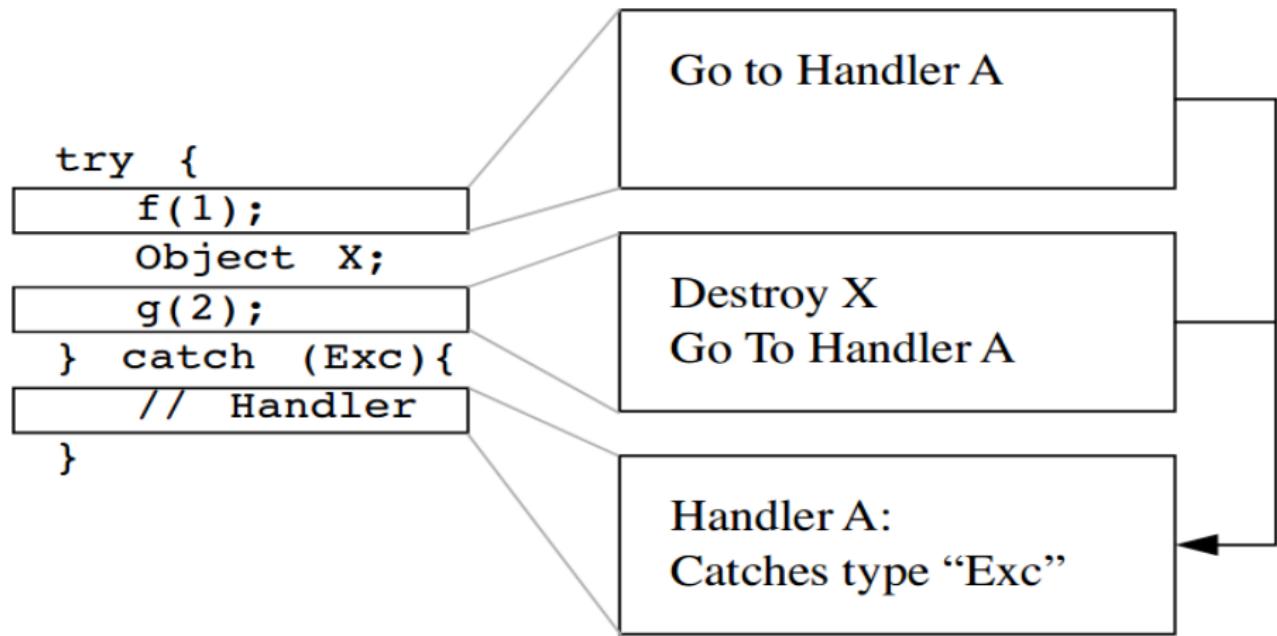


Figure: Table-driven implementation.

Table-driven exceptions

- The compiler translates the throw statement into the calls `__cxa_allocate_exception` and `__cxa_throw`, that allocates the exception and starts the unwinding process
- For every catch statement an entry in the exception table and a cleanup table is created.
- When the exception is thrown, the unwinder goes through the exception table (through a function called personality), until it finds a landing pad.
- If no landing path is found for the exception, `std::terminate` is called.
- If a landing path is found, the unwinder now starts again on the top of the stack.

Table-driven exceptions

- The unwinder goes through the stack a second time and it will call the personality routine to perform a cleanup.
- The personality routine will call the cleanup table. This will call the destructors of the objects allocated.
- The unwinder jumps in the stack into the proper catch statement.
- The memory acquired for the exceptions will be called.

Table-driven exceptions

- Performance penalties
 - The runtime still needs to restore the values of the variables that were declared outside the try block.
 - The destructors of the objects need to be called.
 - The tables have a non-negligible size. They can be stored in the initialized read-only area.
 - No additional cost for the normal path.

Exceptions

```
1 extern int bar(int)
2 //noexcept
3 ;
4
5 int foo() noexcept
6 {
7     int i = 42;
8     int j = 21;
9     return bar(i+j);
10 }
```

Listing 16: Clang 9.0 C++17 -O3

Exceptions

```
1 foo(): # @foo()
2     push rax
3     mov edi, 63
4     call bar(int)
5     pop rcx
6     ret
7     mov rdi, rax
8     call __clang_call_terminate
9 __clang_call_terminate: # @_clang_call_terminate
10    push rax
11    call __cxa_begin_catch
12    call std::terminate()
```

Listing 17: Clang 9.0 C++17 -O3

Exceptions

```
1 extern int bar(int)
2 noexcept
3 ;
4
5 int foo() noexcept
6 {
7     int i = 42;
8     int j = 21;
9     return bar(i+j);
10 }
```

Listing 18: Clang 9.0 C++17 -O3

Exceptions

```
1 foo(): # @foo()
2     mov edi, 63
3     jmp bar(int) # TAILCALL
```

Listing 19: Clang 9.0 C++17 -O3

Exceptions don't have a zero cost!!!



Figure: Exceptions don't have zero cost!

- As seen exceptions don't have a zero cost abstraction. They do have a cost in space and in time, even if they are not used!
- If the code is not using exceptions, make sure that the `-fno-exceptions` is enabled in the compiler.

Excpetions: the call graph is difficult to follow...

- How many execution paths are in the following code?

```
1 String EvaluateSalaryAndReturnName( Employee e )
2 {
3     if( e.Title() == "CEO" || e.Salary() > 100000 )
4     {
5         cout << e.First() << " " << e.Last()
6             << " is overpaid" << endl;
7     }
8     return e.First() + " " + e.Last();
9 }
10
11
12
13
```

Listing 20: From Herb Sutter's Exceptional C++

Exceptions: Basic guarantee, Strong guarantee and noexcept guarantee

- Why we cannot get and object in one operation from a stack? Ex. stack `top()` and `pop()` functions.
 - `void std :: stack < T, Container >:: pop()`
 - `reference std :: stack < T, Container >:: top();`
 - `const_reference std :: stack < T, Container >:: top() const;`
- In node based containers, the `extract` member function has been added ex. no more find erase needed in map. It is also the only way to take a move-only object.

Exceptions

```
1 int main() {
2     std::map<int, std::string> m{{1, "Knuth"}, {2, "Hoare"}, {3,
3 "Stroustrup"}};
4     for(const auto& n : m)
5         std::cout << n.first << " " << n.second << std::endl;
6     auto nh = m.extract(2);
7     nh.key() = 4;
8     m.insert(std::move(nh));
9     for(const auto& n : m)
10        std::cout << n.first << " " << n.second << std::endl;
}
```

Listing 21: Extract from a container

Exceptions: Conclusion

- They are general approach to handle errors
- They are designed with centralized handling in mind.
- They can easily add arbitrary amount of information.
- "Little" cost on normal path.
- Hard to write exception code correctly, specially libraries
- Based on type based, and therefore on RTTI = \downarrow non-deterministic in time, requires dynamic allocation.
- Considerable cost in size
- Banned in many coding style standards (e.g., Google)
- "I can't recommend exceptions for hard real time; doing so is a research problem, which I expect to be solved within the decade [Stroustrup 2004]" .
- Invisible flow

Content

1 Introduction

2 Contracts

3 Error Codes

4 Exceptions

5 Expected

6 Throwing by value

7 Conclusion

8 Questions

Expected: Genesis



Figure: Alexander Alexandrescu.

- In 2012, Alexandrescu gave a talk - Systematic Error Handling in C++ - at C++ and Beyond.
- There he discussed the advantages of a Algebraic data type.

Expected: How do other languages manage errors?

- Rust: It has two type of errors. Recoverable errors and non-recoverable errors.
- $\text{Result} < T, E >$ for recoverable errors and the `panic!` macro that stops execution when the program encounters an unrecoverable error.

Expected: Rust

```
1 enum Result<T, E> {
2     Ok(T),
3     Err(E),
4 }
5
6 fn main() {
7     let f = std::fs::File::open("hello.txt");
8     let f = match f {
9         Ok(file) => file,
10        Err(error) => {
11            panic!("Problem opening the file: {:?}", error)
12        },
13    },
14 };
15 }
```

Listing 22: Pattern matching in Rust

Expected: Rust error handling

```
1 use std::io;
2 use std::io::Read;
3 use std::fs::File;
4
5 fn read_username_from_file() -> Result<String, io::Error> {
6     let mut s = String::new();
7
8     File::open("hello.txt")?.read_to_string(&mut s)?;
9
10    Ok(s)
11 }
12 }
```

Listing 23: Shortcut for errors in Rust

Expected: Do we have something similar in C++?

- No. Yes. Maybe. Nowadays in C++ just `std::optional` exists.
- A `std::expected` class is expected to come in C++23 (first draft in 2014 N4015. Last P0323 in 2019)
- It is a Algebraic data type that returns either the expected value or the error.
- Boost.Outcome is a very similar class designed for such a purpose. Best available option if something like `std::expected` is needed today.

Expected from P0323

```
1
2 template<class T, class E>
3 class expected {
4 public:
5     using value_type = T;
6     using error_type = E;
7     using unexpected_type = unexpected<E>;
8
9 // 4.6, observers
10    constexpr const T* operator->() const;
11    constexpr T* operator->();
12    constexpr const T& operator*() const&;
13    constexpr T& operator*() &;
14    constexpr const T&& operator*() const&&;
15    constexpr T&& operator*() &&;
16    constexpr explicit operator bool() const noexcept;
17    constexpr bool has_value() const noexcept;
18    constexpr const T& value() const&;
19    constexpr T& value() &;
20    constexpr const T&& value() const&&;
21    constexpr T&& value() &&;
22
23 };
```

Expected

```
1  constexpr const E& error() const&;
2  constexpr E& error() &;
3  constexpr const E&& error() const&&;
4  constexpr E&& error() &&;
5  template<class U>
6      constexpr T value_or(U&&) const&;
7  template<class U>
8      constexpr T value_or(U&&) &&;
9
10
11 private:
12     bool has_val; // exposition only
13     union
14     {
15         value_type val; // exposition only
16         unexpected_type unex; // exposition only
17     };
18 }
```

Listing 25: Expected class from P0323

Expected: arithmetic function

```
1
2 enum class arithmetic_errc{
3     divide_by_zero = 1,
4     not_integer_division = 2,
5     integer_divide_overflows = 3,
6 }
7
8 expected<int, arithmetic_errc> safe_divide(int i, int j){
9     if(j == 0)
10         return unexpected(arithmetic_errc::divide_by_zero);
11     if(i == std::numeric_limits<int>::min() && j == -1)
12         return unexpected(arithmetic_errc::integer_divide_overflows);
13     if(i%j != 0)
14         return unexpected(arithmetic_errc::not_integer_division);
15     return i /j;
16 }
```

Listing 26: Ex. from P0323

Expected

```
1
2 expected<double, arithmetic_errc> f( double i, int j, int k){
3     auto q = safe_divide(j,k);
4     if(!q) return unexpected(q.error());
5     return i + q.value(); // i + *q;
6 }
```

Listing 27: Error handling with expected..

Excepted: Conclusion

- Locally to handle errors
- They can easily add arbitrary amount of information (No two exceptions can run on the same time)
- Still tedious to write it correctly
- Cheap, as the value is passed on the stack.
- Deterministics.
- Visible flow.

Expected: Monads

```
1     std::expected<image,fail_reason> get_cute_cat (const image& img
2 ) {
3     auto cropped = crop_to_cat(img);
4     if (!cropped) {
5         return cropped;
6     }
7
8     auto with_tie = add_bow_tie(*cropped);
9     if (!with_tie) {
10        return with_tie;
11    }
12
13    auto with_sparkles = make_eyes_sparkle(*with_tie);
14    if (!with_sparkles) {
15        return with_sparkles;
16    }
17
18    return add_rainbow(make_smaller(*with_sparkles));
19 }
```

Listing 28: Monadic interfaces. From Simon Brand implementation

Expected: Monads

```
1   tl::expected<image,fail_reason> get_cute_cat (const image& img)
2 {
3     return crop_to_cat(img)
4       .and_then(add_bow_tie)
5       .and_then(make_eyes_sparkle)
6       .map(make_smaller)
7       .map(add_rainbow);
8 }
```

Listing 29: Monadic interfaces

Expected: Monads

- map: carries out some operation on the stored object if there is one.
- $tl :: expected < std :: size_t, std :: error_code > s = exp_string.map(std :: string :: size);$
- map_error: carries out some operation on the unexpected object if there is one.
- my_error_code translate_error (std::error_code);
- $tl :: expected < int, my_error_code > s = exp_int.map_error(translate_error);$
- and_then: like map, but for operations which return a tl::expected.
- $tl :: expected < ast, fail_reason > parse(const std :: strings);$
- $tl :: expected < ast, fail_reason > exp_ast = exp_string.and_then(parse);$
- or_else: calls some function if there is no value stored.
- $exp.or_else([] throw std :: runtime_error "ohno";);$

Content

- 1 Introduction
- 2 Contracts
- 3 Error Codes
- 4 Exceptions
- 5 Expected
- 6 Throwing by value
- 7 Conclusion
- 8 Questions

Throwing by value

- Herb Sutter came with another idea...
- Zero-overhead deterministic exceptions: Throwing values: P0709
- “A big appeal to Go using error codes is as a rebellion against the overly complex languages in today’s landscape. We have lost a lot of what makes C so elegant – that you can usually look at any line of code and guess what machine code it translates into.... You’re in a statically typed programming language, and the dynamic nature of exceptions is precisely the reason they suck.” — [Duffy 2016]
- Problem: “Filesystem library functions often provide two overloads, one that throws an exception to report file system errors, and another that sets an `error_code` [N3239]”.
- Different error handling methods are proliferating.
- “... if STL switched to making bugs logic errors, domain errors be contracts... and we could make `bad_alloc` fail fast, we could use the standard STL unmodified throughout Windows.” — Pavel Curtis, private communication between Herb Sutter

Throwing by value

- “This proposal aims to marry the best of exceptions and error codes: to allow a function to declare that it throws values of a statically known type“ P0709
- Throwing such values behaves as-if the function returned `union R; E; +bool`
- On success the function returns the normal return value `R` and on error the function returns the error value type `E`, both in the same return channel including using the same registers. The discriminant can use an unused CPU flag or a register.
- Basically it proposes to bake into the language `std::expected...`



Figure: C++ language baker

Throwing by value

- Throwing std::error_codes in the stack through the new keyword *throws*

```
1
2 string f() throws {
3     if (flip_a_coin())
4         throw arithmetic_error::something;
5     return std::string(xyzzy) + std::string(plover);
6     // any dynamic exception is translated to error
7 }
8 string g() throws { return f() + plugh; }
9 // any dynamic exception is translated to error
10 int main() {
11     try {
12         auto result = g();
13         cout << success, result is: << result;
14     }
15     catch(error err) {
16         // catch by value is fine
17         cout << failed, error is: << err.error();
18     }
19 }
```

Listing 30: Ex. from P0323

Throwing by value

```
1 expected<int, errc> safe_divide(int i, int j) {
2     if (j == 0)
3         return unexpected(arithmetic_errc::divide_by_zero);
4     if (i == INT_MIN && j == -1)
5         return unexpected(arithmetic_errc::integer_divide_overflows);
6     if (i % j != 0)
7         return unexpected(arithmetic_errc::not_integer_division);
8     return i / j;
9 }
10
11
12 expected<double, errc> caller(double i, double j, double k) {
13     auto q = safe_divide(j, k);
14     if (q) return i + *q;
15     else return q;
16 }
```

Listing 31: Ex. from P0323R3

Throwing by value

```
1 int safe_divide(int i, int j) throws {
2     if (j == 0)
3         throw arithmetic_errc::divide_by_zero;
4     if (i == INT_MIN && j == -1)
5         throw arithmetic_errc::integer_divide_overflows;
6     if (i % j != 0)
7         throw arithmetic_errc::not_integer_division;
8     return i / j;
9 }
10
11 double caller(double i, double j, double k) throws {
12     return i + safe_divide(j, k);
13 }
```

Listing 32: Ex. from P0709R0

Throwing by value: Heap exhaustion debate

- In favour
 - It is an explicit fallible request in source code, unlike stack exhaustion
 - Some (possibly lots) of code is correct today for heap exhaustion even without coding for it specially. Calling new asking for too much data.
- Against
 - Testing is much more difficult. Specially for heap exhaustion.
 - Recovery requires special care. Many programs that believe that they can recover, they are buggy.
 - Treating heap exhaustion the same as all other errors causes a overhead in the cases that don't need it.
 - Many programs nowadays fail on heap allocation failure (e.g., google)



Figure: Pros and Cons

Throwing by value

- In any case this proposal does not have a target year (e.g., C++23, C++26...)



Figure: Definitely not in C++20..

Content

1 Introduction

2 Contracts

3 Error Codes

4 Exceptions

5 Expected

6 Throwing by value

7 Conclusion

8 Questions

Conclusion

- There co-exists nowadays 3 methods for error handling : error codes, exceptions and outcome/expected algebraic data types.
- A possible 4th method may be baked into the language
- If a new project is started, a outcome/expected approach is the most recommendable.

Content

1 Introduction

2 Contracts

3 Error Codes

4 Exceptions

5 Expected

6 Throwing by value

7 Conclusion

8 Questions

Questions?

Any question?